

# index

March 21, 2022

## 1 Regression Trees and Model Optimization - Lab

### 1.1 Introduction

In this lab, we'll see how to apply regression analysis using CART trees while making use of some hyperparameter tuning to improve our model.

### 1.2 Objectives

In this lab you will:

- Perform the full process of cleaning data, tuning hyperparameters, creating visualizations, and evaluating decision tree models
- Determine the optimal hyperparameters for a decision tree model and evaluate the performance of decision tree models

### 1.3 Ames Housing dataset

The dataset is available in the file 'ames.csv'.

- Import the dataset and examine its dimensions:

```
[6]: # Import necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use('ggplot')
%matplotlib inline

# Load the Ames housing dataset
data = pd.read_csv("ames.csv")

# Print the dimensions of data
data.shape

# Check out the info for the dataframe
data.info()

# Show the first 5 rows
```

```
data.head()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1460 entries, 0 to 1459
```

```
Data columns (total 81 columns):
```

#	Column	Non-Null Count	Dtype
0	Id	1460 non-null	int64
1	MSSubClass	1460 non-null	int64
2	MSZoning	1460 non-null	object
3	LotFrontage	1201 non-null	float64
4	LotArea	1460 non-null	int64
5	Street	1460 non-null	object
6	Alley	91 non-null	object
7	LotShape	1460 non-null	object
8	LandContour	1460 non-null	object
9	Utilities	1460 non-null	object
10	LotConfig	1460 non-null	object
11	LandSlope	1460 non-null	object
12	Neighborhood	1460 non-null	object
13	Condition1	1460 non-null	object
14	Condition2	1460 non-null	object
15	BldgType	1460 non-null	object
16	HouseStyle	1460 non-null	object
17	OverallQual	1460 non-null	int64
18	OverallCond	1460 non-null	int64
19	YearBuilt	1460 non-null	int64
20	YearRemodAdd	1460 non-null	int64
21	RoofStyle	1460 non-null	object
22	RoofMatl	1460 non-null	object
23	Exterior1st	1460 non-null	object
24	Exterior2nd	1460 non-null	object
25	MasVnrType	1452 non-null	object
26	MasVnrArea	1452 non-null	float64
27	ExterQual	1460 non-null	object
28	ExterCond	1460 non-null	object
29	Foundation	1460 non-null	object
30	BsmtQual	1423 non-null	object
31	BsmtCond	1423 non-null	object
32	BsmtExposure	1422 non-null	object
33	BsmtFinType1	1423 non-null	object
34	BsmtFinSF1	1460 non-null	int64
35	BsmtFinType2	1422 non-null	object
36	BsmtFinSF2	1460 non-null	int64
37	BsmtUnfSF	1460 non-null	int64
38	TotalBsmtSF	1460 non-null	int64
39	Heating	1460 non-null	object

```

40 HeatingQC      1460 non-null object
41 CentralAir     1460 non-null object
42 Electrical     1459 non-null object
43 1stFlrSF       1460 non-null int64
44 2ndFlrSF       1460 non-null int64
45 LowQualFinSF   1460 non-null int64
46 GrLivArea      1460 non-null int64
47 BsmtFullBath   1460 non-null int64
48 BsmtHalfBath   1460 non-null int64
49 FullBath       1460 non-null int64
50 HalfBath       1460 non-null int64
51 BedroomAbvGr  1460 non-null int64
52 KitchenAbvGr  1460 non-null int64
53 KitchenQual    1460 non-null object
54 TotRmsAbvGrd  1460 non-null int64
55 Functional     1460 non-null object
56 Fireplaces     1460 non-null int64
57 FireplaceQu    770 non-null object
58 GarageType     1379 non-null object
59 GarageYrBlt    1379 non-null float64
60 GarageFinish   1379 non-null object
61 GarageCars     1460 non-null int64
62 GarageArea     1460 non-null int64
63 GarageQual     1379 non-null object
64 GarageCond     1379 non-null object
65 PavedDrive     1460 non-null object
66 WoodDeckSF     1460 non-null int64
67 OpenPorchSF    1460 non-null int64
68 EnclosedPorch  1460 non-null int64
69 3SsnPorch      1460 non-null int64
70 ScreenPorch    1460 non-null int64
71 PoolArea       1460 non-null int64
72 PoolQC        7 non-null object
73 Fence          281 non-null object
74 MiscFeature    54 non-null object
75 MiscVal        1460 non-null int64
76 MoSold         1460 non-null int64
77 YrSold         1460 non-null int64
78 SaleType       1460 non-null object
79 SaleCondition  1460 non-null object
80 SalePrice      1460 non-null int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

```

[6]:   Id  MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape \
0    1          60      RL         65.0    8450   Pave   NaN      Reg
1    2          20      RL         80.0    9600   Pave   NaN      Reg

```

2	3	60	RL	68.0	11250	Pave	NaN	IR1
3	4	70	RL	60.0	9550	Pave	NaN	IR1
4	5	60	RL	84.0	14260	Pave	NaN	IR1

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	\
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	5	
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0	9	
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0	12	

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

## 1.4 Identify features and target data

In this lab, we will use using 3 predictive continuous features:

### Features

- LotArea: Lot size in square feet
- 1stFlrSF: Size of first floor in square feet
- GrLivArea: Above grade (ground) living area square feet

### Target

- SalePrice', the sale price of the home, in dollars
- Create DataFrames for the features and the target variable as shown above
- Inspect the contents of both the features and the target variable

```
[7]: # Features and target data
target = data["SalePrice"]
features = data[["LotArea", "1stFlrSF", "GrLivArea"]]
```

## 1.5 Inspect correlations




- Use scatter plots to show the correlation between the chosen features and the target variable
- Comment on each scatter plot

```
[12]: # Your code here

fig, axes = plt.subplots(nrows = 3, ncols = 1, figsize = (5,15))
```

```
fig.subplots_adjust(hspace=0.4, wspace=0.25)

for i,item in enumerate(features.columns):
    ax = axes[i]
    sns.scatterplot(x = features[item], y = target, ax = ax, label = item,
                    color = "tab:blue")
```



## 1.6 Create evaluation metrics

- Import `r2_score` and `mean_squared_error` from `sklearn.metrics`
- Create a function `performance(true, predicted)` to calculate and return both the R-squared score and Root Mean Squared Error (RMSE) for two equal-sized arrays for the given true and predicted values
  - Depending on your version of sklearn, in order to get the RMSE score you will need to either set `squared=False` or you will need to take the square root of the output of the `mean_squared_error` function - check out [the documentation](#) or this helpful and related [StackOverflow post](#)
  - The benefit of calculating RMSE instead of the Mean Squared Error (MSE) is that RMSE is in the same units as the target - here, this means that RMSE will be in dollars, calculating how far off in dollars our predictions are away from the actual prices for homes, on average

```
[26]: # Import metrics
from sklearn.metrics import r2_score, mean_squared_error

# Define the function
def performance(y_true, y_predict):
    """
    Calculates and returns the two performance scores between
    true and predicted values - first R-Squared, then RMSE
    """

    # Calculate the r2 score between 'y_true' and 'y_predict'
    R2 = r2_score(y_true, y_predict)
    # Calculate the root mean squared error between 'y_true' and 'y_predict'
    RMSE = mean_squared_error(y_true, y_predict, squared=False)
    # Return the score

    return R2, RMSE

# Test the function
score = performance([3, -0.5, 2, 7, 4.2], [2.5, 0.0, 2.1, 7.8, 5.3])
score

# [0.9228556485355649, 0.6870225614927066]
```

```
[26]: (0.9228556485355649, 0.6870225614927066)
```

## 1.7 Split the data into training and test sets

- Split features and target datasets into training/test data (80/20)

- For reproducibility, use `random_state=42`

```
[27]: from sklearn.model_selection import train_test_split

# Split the data into training and test subsets
x_train, x_test, y_train, y_test = train_test_split(features, target,
                                                    test_size = 0.2,
                                                    random_state = 42)
```

## 1.8 Grow a vanilla regression tree

- Import the `DecisionTreeRegressor` class
- Run a baseline model for later comparison using the datasets created above
- Generate predictions for test dataset and calculate the performance measures using the function created above
- Use `random_state=45` for tree instance
- Record your observations

```
[28]: # Import DecisionTreeRegressor

from sklearn.tree import DecisionTreeRegressor

# Instantiate DecisionTreeRegressor
# Set random_state=45
regressor = DecisionTreeRegressor(random_state = 45)

# Fit the model to training data
regressor.fit(x_train, y_train)

# Make predictions on the test data
y_pred = regressor.predict(x_test)

# Calculate performance using the performance() function
score = performance(y_test, y_pred)
score

# [0.5961521990414137, 55656.48543887347] - R2, RMSE
```

```
[28]: (0.5961521990414137, 55656.48543887347)
```

## 1.9 Hyperparameter tuning (I)

- Find the best tree depth using depth range: 1-30
- Run the regressor repeatedly in a `for` loop for each depth value
- Use `random_state=45` for reproducibility
- Calculate RMSE and r-squared for each run
- Plot both performance measures for all runs



- Comment on the output

```
[38]: # Your code here

depth = range(1, 31)

R2 = []
RMSE = []

for i in depth:


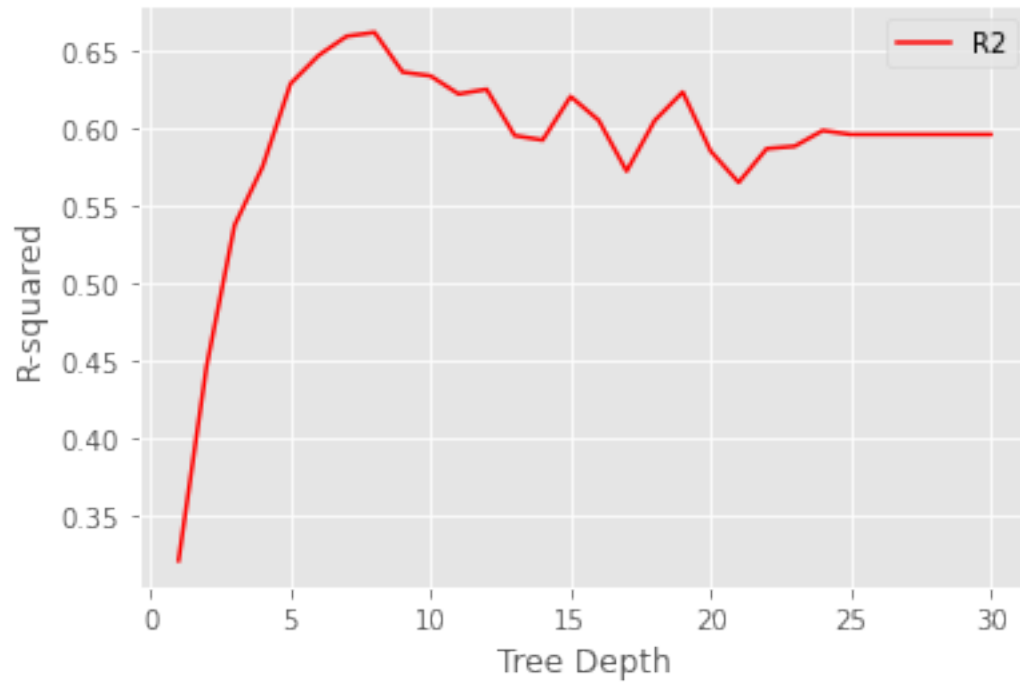
    regressor = DecisionTreeRegressor(random_state = 45, max_depth = i)

    # Fit the model to training data
    regressor.fit(x_train, y_train)

    # Make predictions on the test data
    y_pred = regressor.predict(x_test)

    # Calculate performance using the performance() function
    score = performance(y_test, y_pred)
    R2.append(score[0])
    RMSE.append(score[1])

sns.lineplot(x = depth, y = R2, label = "R2", color = "r");
plt.xlabel('Tree Depth')
plt.ylabel('R-squared')
plt.show()
sns.lineplot(x = depth, y = RMSE, label = "RMSE", color = "blue")
plt.xlabel('Tree Depth')
plt.ylabel('RMSE')
plt.show()
```



```
[40]: # We can see that at tree depth that is equal to 7, we have minimum value for
      # RMSE and maximum value for R2. So max_depth = 7 is the optimum value
```

## 1.10 Hyperparameter tuning (II)

- Repeat the above process for `min_samples_split`
- Use a range of values from 2-10 for this hyperparameter
- Use `random_state=45` for reproducibility
- Visualize the output and comment on results as above

```
[39]: # Your code here

min_samples_split = range(2, 11)

R2 = []
RMSE = []

for i in min_samples_split:

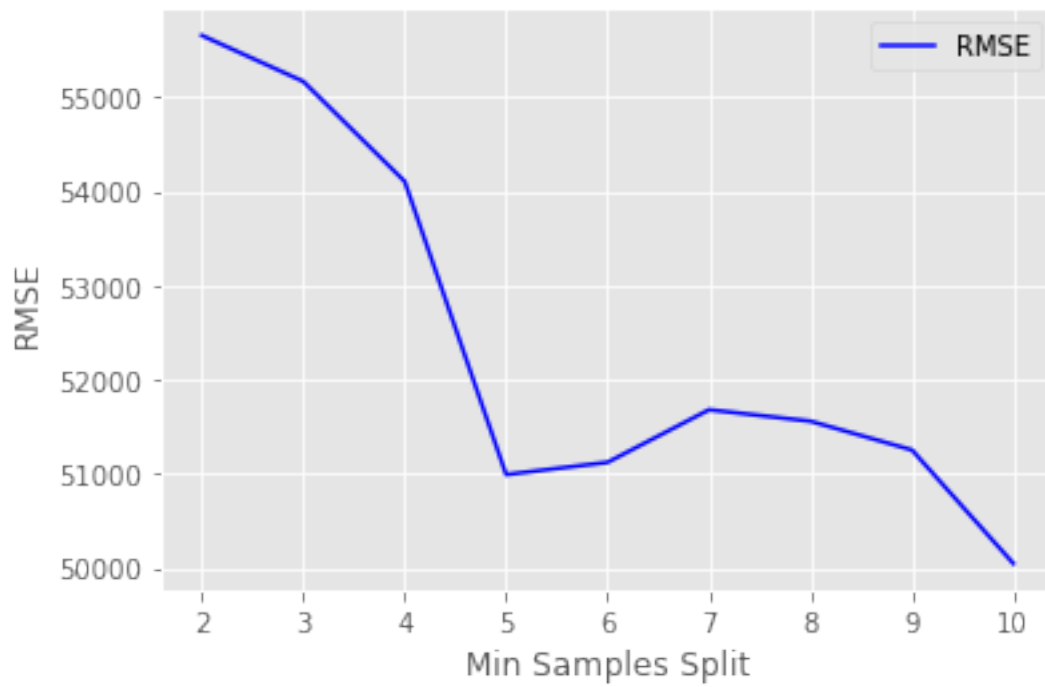
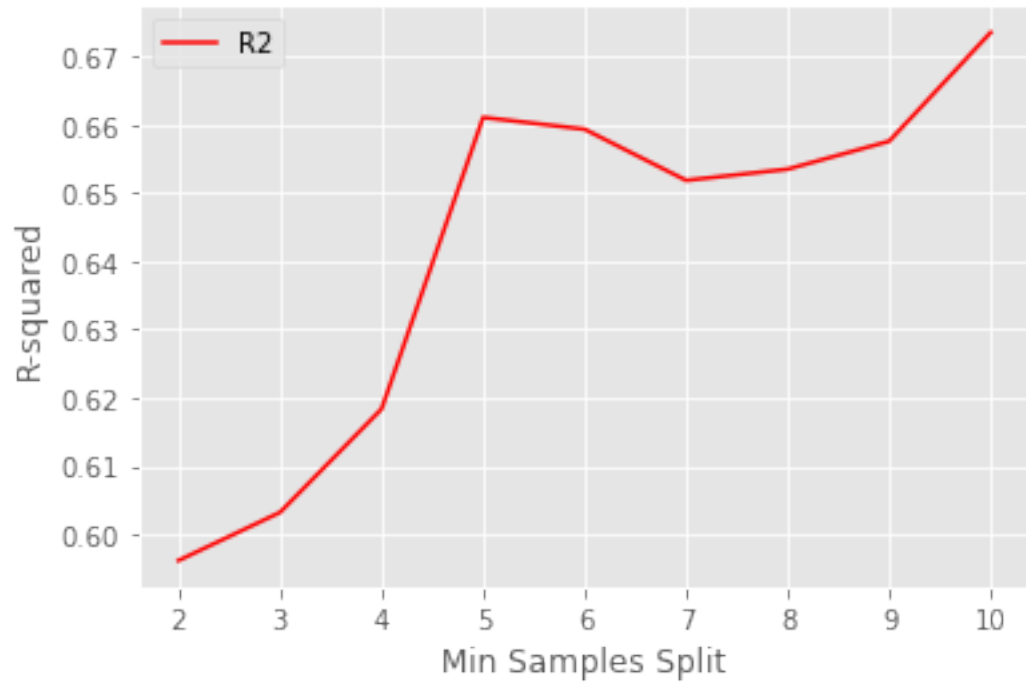
    regressor = DecisionTreeRegressor(random_state = 45, min_samples_split=i)

    # Fit the model to training data
    regressor.fit(x_train, y_train)

    # Make predictions on the test data
    y_pred = regressor.predict(x_test)

    # Calculate performance using the performance() function
    score = performance(y_test, y_pred)
    R2.append(score[0])
    RMSE.append(score[1])

sns.lineplot(x = min_samples_split, y = R2, label = "R2", color = "r");
plt.xlabel('Min Samples Split')
plt.ylabel('R-squared')
plt.show()
sns.lineplot(x = min_samples_split, y = RMSE, label = "RMSE", color = "blue")
plt.xlabel('Min Samples Split')
plt.ylabel('RMSE')
plt.show()
```



## 2 Run the *optimized* model

- Use the best values for `max_depth` and `min_samples_split` found in previous runs and run an optimized model with these values
- Calculate the performance and comment on the output

```
[44]: # Your code here

regressor = DecisionTreeRegressor(random_state = 45, min_samples_split=10,
                                   max_depth = 7)

# Fit the model to training data
regressor.fit(x_train, y_train)

# Make predictions on the test data
y_pred = regressor.predict(x_test)

# Calculate performance using the performance() function
score = performance(y_test, y_pred)
print("R2    :", score[0])
print("RMSE :", score[1])
```

```
R2    : 0.6772046046188305
RMSE  : 49758.87841130982
```

### 2.1 Level up (Optional)

- How about bringing in some more features from the original dataset which may be good predictors?
- Also, try tuning more hyperparameters like `max_features` to find a more optimal version of the model

```
[ ]: # Your code here
```

### 2.2 Summary

In this lab, we looked at applying a decision-tree-based regression analysis on the Ames Housing dataset. We saw how to train various models to find the optimal values for hyperparameters.